



Vidyavardhini's College of Engineering & Technology

Department of Computer Engineering

Academic Year : 2023-24

Experiment No. 2

Analyze the Titanic Survival Dataset and apply appropriate regression technique

Date of Performance:

Date of Submission:



Aim: Analyze the Titanic Survival Dataset and apply appropriate Regression Technique.

Objective: Able to perform various feature engineering tasks, apply logistic regression on the given dataset and maximize the accuracy.

Theory:

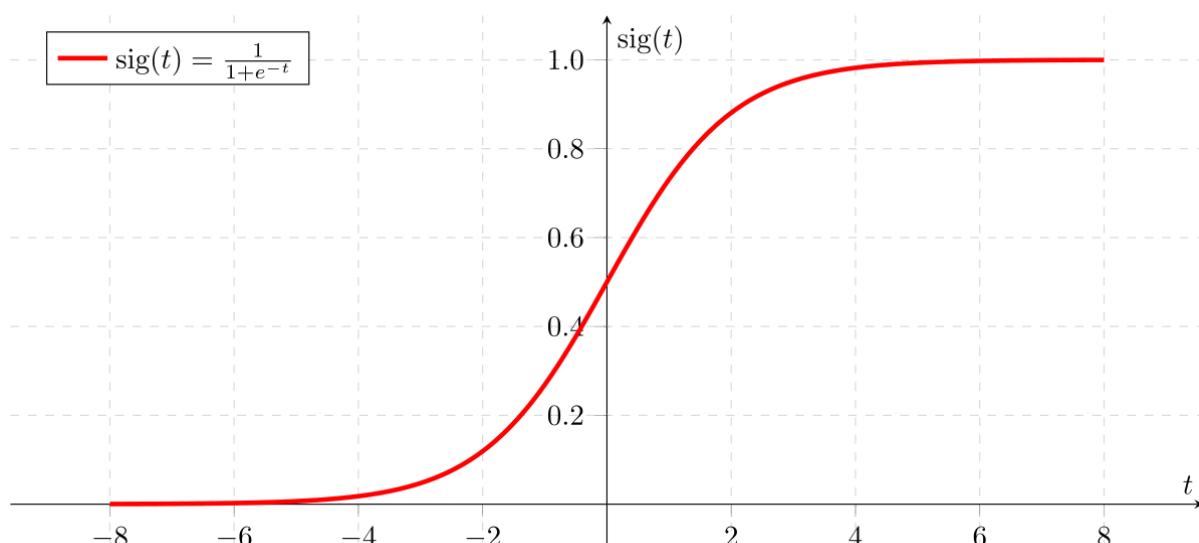
Logistic Regression was used in the biological sciences in early twentieth century. It was then used in many social science applications. Logistic Regression is used when the dependent variable(target) is categorical and is binary in nature. In order to perform binary classification the logistic regression techniques makes use of Sigmoid function.

For example,

To predict whether an email is spam (1) or (0)

Whether the tumor is malignant (1) or not (0)

Consider a scenario where we need to classify whether an email is spam or not. If we use linear regression for this problem, there is a need for setting up a threshold based on which classification can be done. Say if the actual class is malignant, predicted continuous value 0.4 and the threshold value is 0.5, the data point will be classified as not malignant which can lead to serious consequence in real time.





Vidyavardhini's College of Engineering & Technology

Department of Computer Engineering

Academic Year : 2023-24

From this example, it can be inferred that linear regression is not suitable for classification problem. Linear regression is unbounded, and this brings logistic regression into picture. Their value strictly ranges from 0 to 1.

Dataset:

The sinking of the Titanic is one of the most infamous shipwrecks in history.

On April 15, 1912, during her maiden voyage, the widely considered “unsinkable” RMS Titanic sank after colliding with an iceberg. Unfortunately, there weren’t enough lifeboats for everyone onboard, resulting in the death of 1502 out of 2224 passengers and crew.

While there was some element of luck involved in surviving, it seems some groups of people were more likely to survive than others.

In this challenge, we ask you to build a predictive model that answers the question: “what sorts of people were more likely to survive?” using passenger data (ie name, age, gender, socio-economic class, etc).

Variable	Definition	Key
survival	Survival	0 = No, 1 = Yes
pclass	Ticket class	1 = 1st, 2 = 2nd, 3 = 3rd
sex	Sex	
Age	Age in years	
sibsp	# of siblings / spouses aboard the Titanic	
parch	# of parents / children aboard the Titanic	
ticket	Ticket number	
fare	Passenger fare	
cabin	Cabin number	



Vidyavardhini's College of Engineering & Technology

Department of Computer Engineering

Academic Year : 2023-24

embarke
d

Port of Embarkation

C = Cherbourg, Q = Queenstown, S =
Southampton

Variable Notes

pclass: A proxy for socio-economic status (SES)

1st = Upper, 2nd = Middle, 3rd = Lower

age: Age is fractional if less than 1. If the age is estimated, is it in the form of xx.5

sibsp: The dataset defines family relations in this way...,

Sibling = brother, sister, stepbrother, stepsister

Spouse = husband, wife (mistresses and fiancés were ignored)

parch: The dataset defines family relations in this way...

Parent = mother, father

Child = daughter, son, stepdaughter, stepson

Some children travelled only with a nanny, therefore parch=0 for them.

Code:

```

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix
import warnings
warnings.filterwarnings('ignore')

df = pd.read_csv('train.csv')
df.head()

```

PassengerId Survived Pclass Name Sex Age SibSp Parch Ticket Fare

0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500	
1	2	1	1	Cumings, Mrs. John Bradley (Florence D----	female	38.0	1	0	PC 17599	71.2833	

```
df.shape
```

```
(891, 12)
```

```
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   PassengerId 891 non-null    int64  
 1   Survived     891 non-null    int64  
 2   Pclass       891 non-null    int64  
 3   Name         891 non-null    object 
 4   Sex          891 non-null    object 
 5   Age          714 non-null    float64 
 6   SibSp        891 non-null    int64  
 7   Parch        891 non-null    int64  
 8   Ticket       891 non-null    object 
 9   Fare          891 non-null    float64 
 10  Cabin         204 non-null    object 
 11  Embarked     889 non-null    object 
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB

```

```
df.isnull().sum()
```

```

PassengerId      0
Survived        0
Pclass          0
Name            0
Sex             0
Age           177
SibSp          0
Parch          0
Ticket         0
Fare           0
Cabin         687
Embarked       2
dtype: int64

```

```
df = df.drop(columns='Cabin', axis=1)
```

```
df['Age'].fillna(df['Age'].mean(), inplace=True)
```

```
print(df['Embarked'].mode())
```

<https://colab.research.google.com/drive/1x64URkDFSzfdUofQ4i4tr-tNO9VKW2Ze#scrollTo=DVLxgIk9ISqM&printMode=true>

```
print(df['Embarked'].mode()[0])
```

```
0    S
dtype: object
```

```
print(df['Embarked'].mode()[0])
```

```
S
```

```
df['Embarked'].fillna(df['Embarked'].mode()[0], inplace=True)
```

```
df.isnull().sum()
```

	PassengerId	Survived	Pclass	Age	SibSp	Parch
count	891.000000	891.000000	891.000000	891.000000	891.000000	891.000000
mean	446.000000	0.383838	2.308642	29.699118	0.523008	0.381594
std	257.353842	0.486592	0.836071	13.002015	1.102743	0.806057
min	1.000000	0.000000	1.000000	0.420000	0.000000	0.000000
25%	223.500000	0.000000	2.000000	22.000000	0.000000	0.000000
50%	446.000000	0.000000	3.000000	29.699118	0.000000	0.000000
75%	668.500000	1.000000	3.000000	35.000000	1.000000	0.000000
max	891.000000	1.000000	3.000000	80.000000	8.000000	6.000000

```
dtype: int64
```

Data Analysis

```
df.describe()
```

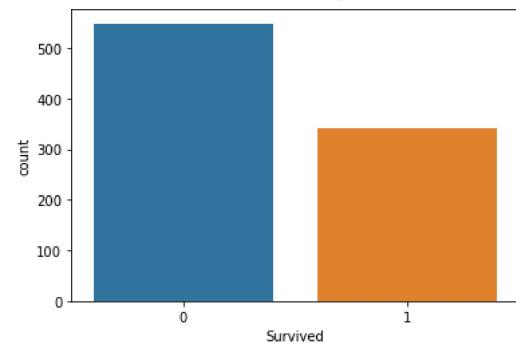
	PassengerId	Survived	Pclass	Age	SibSp	Parch	
count	891.000000	891.000000	891.000000	891.000000	891.000000	891.000000	891.000000
mean	446.000000	0.383838	2.308642	29.699118	0.523008	0.381594	32.2
std	257.353842	0.486592	0.836071	13.002015	1.102743	0.806057	49.6
min	1.000000	0.000000	1.000000	0.420000	0.000000	0.000000	0.0
25%	223.500000	0.000000	2.000000	22.000000	0.000000	0.000000	7.9
50%	446.000000	0.000000	3.000000	29.699118	0.000000	0.000000	14.4
75%	668.500000	1.000000	3.000000	35.000000	1.000000	0.000000	31.0
max	891.000000	1.000000	3.000000	80.000000	8.000000	6.000000	512.3

```
df['Survived'].value_counts()
```

```
0    549
1    342
Name: Survived, dtype: int64
```

```
sns.countplot(x='Survived', data=df)
```

```
<AxesSubplot:xlabel='Survived', ylabel='count'>
```

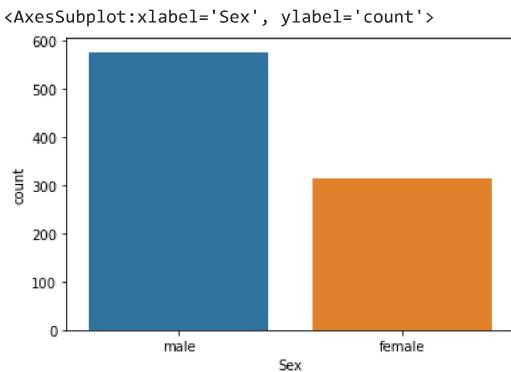


```
df['Sex'].value_counts()
```

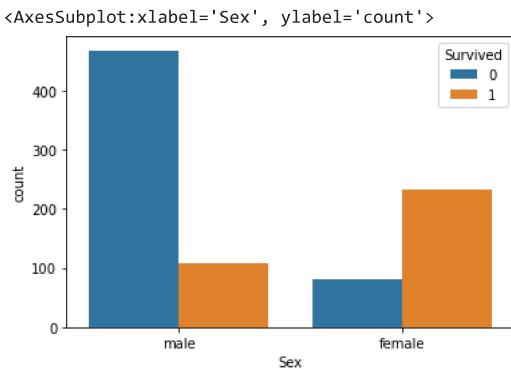
```
male    577
female  314
Name: Sex, dtype: int64
```

```
Name: Sex, dtype: int64
```

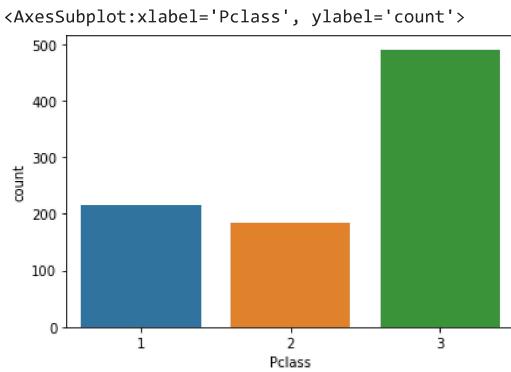
```
sns.countplot(x='Sex', data=df)
```



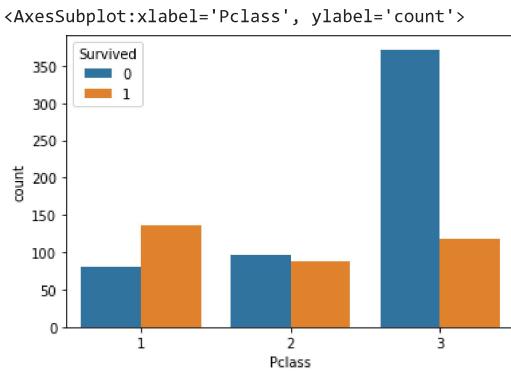
```
sns.countplot(x='Sex', hue='Survived', data=df)
```



```
sns.countplot(x='Pclass', data=df)
```



```
sns.countplot(x='Pclass', hue='Survived', data=df)
```



```
df['Sex'].value_counts()
```

male	577
female	314
Name:	Sex, dtype: int64

```
df['Embarked'].value_counts()
```

S	646
C	168
Q	77
Name:	Embarked, dtype: int64

```
df.replace({'Sex':{'male':0,'female':1}, 'Embarked':{'S':0,'C':1,'Q':2}}, inplace=True)
```

```
df.head()
```

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	F
0	1	0	3	Braund, Mr. Owen Harris	0	22.0	1	0	A/5 21171	7.2
1	2	1	1	Cumings, Mrs. John Bradley (Florence Douce)	1	38.0	1	0	PC 17599	71.2

```
X = df.drop(columns = ['PassengerId','Name','Ticket','Survived'],axis=1)
```

```
Y = df['Survived']
```

```
X.head(3)
```

	Pclass	Sex	Age	SibSp	Parch	Fare	Embarked
0	3	0	22.0	1	0	7.2500	0
1	1	1	38.0	1	0	71.2833	1
2	3	1	26.0	0	0	7.9250	0

```
Y.head(3)
```

0	0
1	1
2	1
Name:	Survived, dtype: int64

```
X_train, X_test, Y_train, Y_test = train_test_split(X,Y, test_size=0.2, random_state=2)
```

```
model = LogisticRegression()
```

```
model.fit(X_train, Y_train)
```

LogisticRegression
LogisticRegression()

```
X_train_prediction = model.predict(X_train)
```

```
X_train_prediction
```

array([0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0,	
1, 0, 0, 1, 0, 0, 1, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0,	
1, 0, 0, 1, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0,	
0, 1, 1, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0,	
1, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0,	
0, 1, 1, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0,	
1, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0,	
0, 1, 1, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0,	
1, 1, 1, 1, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,	
1, 0,	
0, 0, 1, 0, 1, 0, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0,	
0, 1, 1, 0, 1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0,	

```
training_data_accuracy = accuracy_score(Y_train, X_train_prediction)
print('Accuracy score of training data : ', training_data_accuracy)
```

Accuracy score of training data : 0.8075842696629213

```
# accuracy on test data  
X_test_prediction = model.predict(X_test)
```

x_test_prediction

```
test_data_accuracy = accuracy_score(Y_test, X_test_prediction)
print('Accuracy score of test data : ', test_data_accuracy)
```

Accuracy score of test data : 0.7821229050279329

```
print("Confusion matrix :-")
```

```
sns.heatmap(confusion_matrix(Y_test, X_test_prediction), annot=True)
```

Confusion matrix :-
<AxesSubplot:>



```
from sklearn.metrics import classification_report  
print(classification_report(Y_test, X_test_prediction))
```

precision	recall	f1-score	support
0.75	0.91	0.82	100
0.84	0.62	0.72	79

accuracy		0.78	179
macro avg	0.80	0.77	179
weighted avg	0.79	0.78	179



Conclusion:

- 1) The features chosen to develop the model for determining the survival of a passenger are:
 - a. pclass (Passenger Class): Passenger with higher class can illustrate that they have high social political status indicating high chance of survival.
 - b. age (Age): Age can be a critical factor as older people and children can have a low chance of surviving.
 - c. sibsp (Number of Siblings/Spouses Aboard): The number of siblings can also be a deciding factor as more siblings may indicate higher chance of surviving.
 - d. parch (Number of Parents/Children Aboard): This can also be a factor as more parents or children can help in survival.
 - e. Fare: Fare can also be a deciding factor as higher fare may indicate high class travel and thus more chance of surviving.
 - f. Sex (Gender, Male): Gender can also be a critical factor in deciding the chance of surviving.
 - g. Embarked: The Port of Embarkation can also affect the chance of survival depending on location.
- 2) The training accuracy of 80.25% and test accuracy of 78.21% indicates that the model's predictions are similar with the actual outcomes in the database. We also calculated precision, recall, and F1-score metrics to provide additional insights.